Imperial College London

May 5th, 2020

Convolutional Neural Networks

Olivier Dubrule/Navjot Kukreja

1

Imperial College London

Objectives of the Day

- Understand what Convolutional Neural Networks (CNNs) are
- Be at ease with calculations associated with the CNN parameters
- Get familiar with classical CNN structures on well-known examples

2

Imperial College London

Convolutional Neural Networks

1. Convolutional Neural Networks
2. Pooling and Fully Connected Layers
3. Examples

3

Imperial College London

Computer Vision Topics

Image Classification

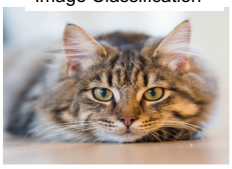

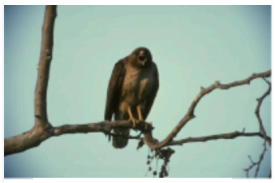

Image Segmentation

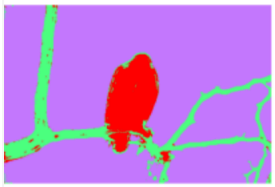

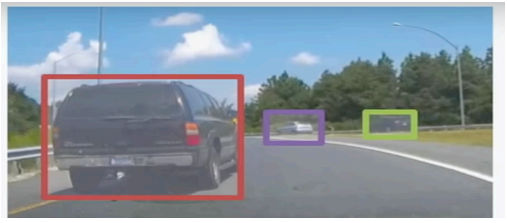

Object Localization (1 object) or Detection (several objects)

4

Imperial College London

MNIST

Extract from the MNIST dataset

MNIST: Modified National Institute of Standards and Technology database

5

Imperial College London

What are Convolutional Neural Networks?

FeedForward NN Approach on MNIST:
Spatial Organization is Lost

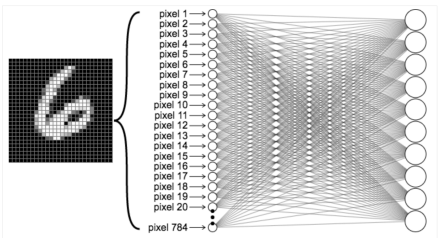

A CNN will take into account the spatial organization of the dataset

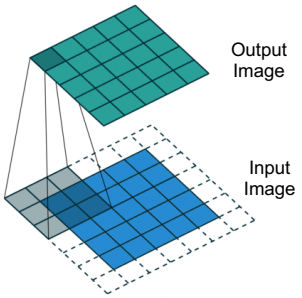

The grey 3x3 moving window is the convolutional filter

6

Imperial College London

What is the Convolution Operation?

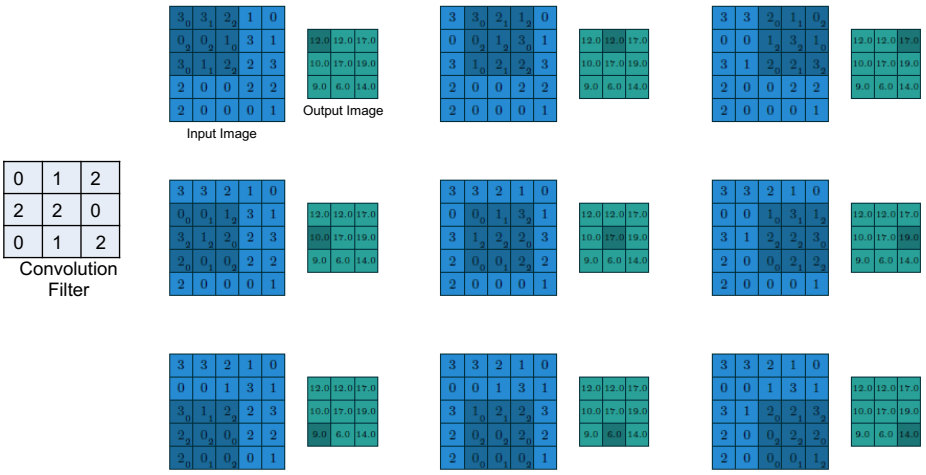

From Dumoulin & Visin, 2018

7

Imperial College London

Exercise: Calculate the Output Image

Input Image

4	2	3	3
2	1	2	5
3	2	0	3
1	2	3	1

*

Convolution Filter (or Kernel)

2	-3	2
1	0	1
2	-3	3

=

Output Image

12	20
13	9

8

Imperial College London

Filters (or Kernels) for Edge Detection

Filter for Detecting Vertical Edges

Input Image

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

* *Convolution Operator*

Filter (or Kernel)

1	0	-1
1	0	-1
1	0	-1

=

Output Image

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

Inspired from Andrew Ng

9

Imperial College London

Which Filter should I use? (1)

Blur

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Horizontal Edges

-1	-1	-1
2	2	2
-1	-1	-1

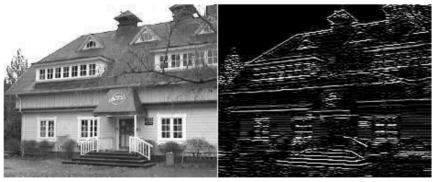

https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/

10

Imperial College London

Which Filter should I use? (2)

Edges

-1	-1	-1
-1	8	-1
-1	-1	-1

Sobel Filters

-1	-2	-1
0	0	0
1	2	1

Horizontal

-1	0	1
-2	0	2
-1	0	1

Vertical

https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/

11

Imperial College London

Idea Behind Convolutional Neural Networks:

Parametrize the Filter and Optimize its Coefficients Based on the Objective!

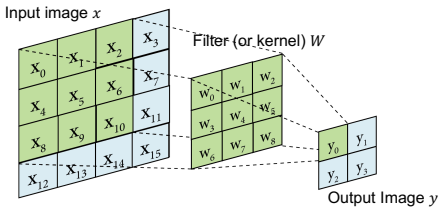

$$W * x = y$$

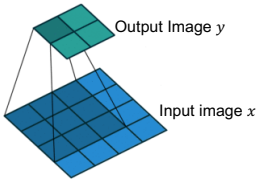

$$y_0 = w_0x_0+w_1x_1+w_2x_2+w_3x_4+w_4x_5+w_5x_6+w_6x_8+w_7x_9+w_8x_{10}$$

12

Imperial College London

Making Each Kernel Operation Non-Linear

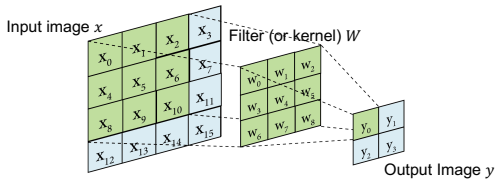

$$y_0 = w_0x_0+w_1x_1+w_2x_2+w_3x_4+w_4x_5+w_5x_6+w_6x_8+w_7x_9+w_8x_{10}$$

A bias term can be added

$$y_0 = w_0x_0+w_1x_1+w_2x_2+w_3x_4+w_4x_5+w_5x_6+w_6x_8+w_7x_9+w_8x_{10} + \boldsymbol{b}$$

And a non-linear transformation is applied by an activation function g

$$y_0 = \boldsymbol{g}(w_0x_0+w_1x_1+w_2x_2+w_3x_4+w_4x_5+w_5x_6+w_6x_8+w_7x_9+w_8x_{10} + \boldsymbol{b})$$

13

Imperial College London

Linear Convolution Operation….

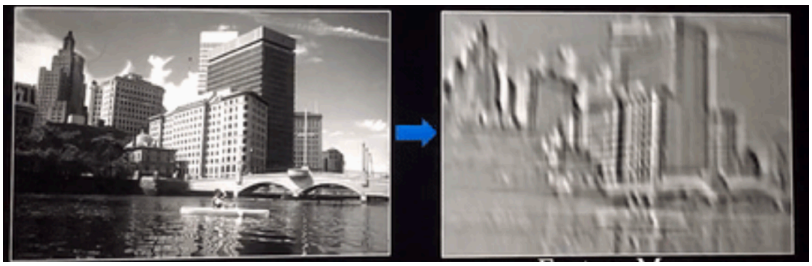

https://ujwlkarn.files.wordpress.com/2016/08/giphy.gif?w=480&zoom=2

14

Imperial College London

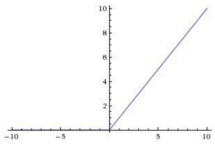

.....Followed by ReLU Activation

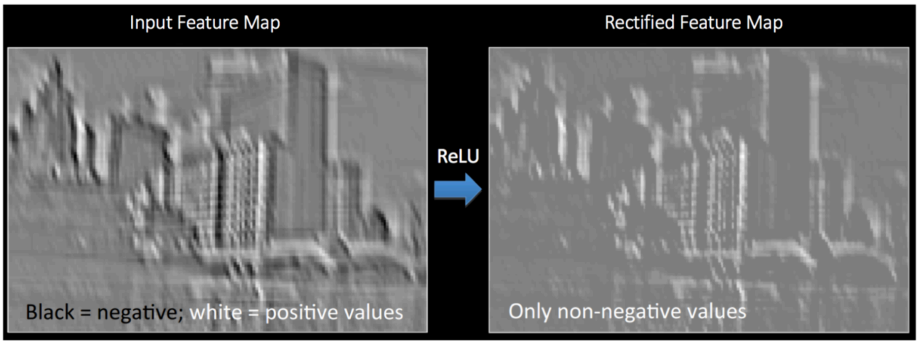

https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/

15

Imperial College London

Key Parameters of Convolutional Neural Networks

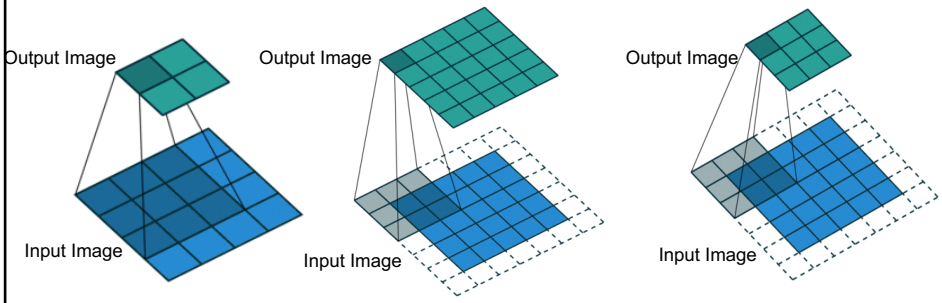

No padding, Stride 1 **Padding 1, Stride 1** **Padding 1, Stride 2**

16

Imperial College London

The Two-Dimensional Convolution Parameters

Parameters

Input Image Size: $n \times n$

Filter (or Kernel) Size: f $\times f$

Padding: p

Stride : s

Output Image Size

$$\left(\frac{n+2p-f}{s}+1\right) \times \left(\frac{n+2p-f}{s}+1\right)$$

Input Image

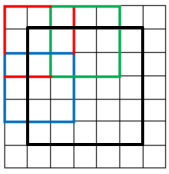

Output Image

3x3!

$n = 5\ p = 1\ f = 3\ s = 2$

17

Imperial College London

The Two-Dimensional Convolution Parameters

Parameters

Image Size: $n \times n$

Filter (or Kernel) Size: f $\times f$

Padding: p

Stride : s

Output Image Size

$$\left(\frac{n+2p-f}{s}+1\right) \times \left(\frac{n+2p-f}{s}+1\right)$$

Input Image

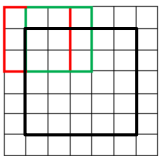

Output Image

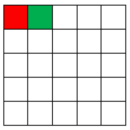

5x5!

$n = 5\ p = 1\ f = 3\ s = 1$

18

Imperial College London

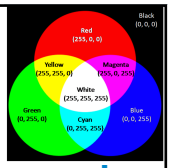

If the Two-Dimensional image has 3 RGB channels…

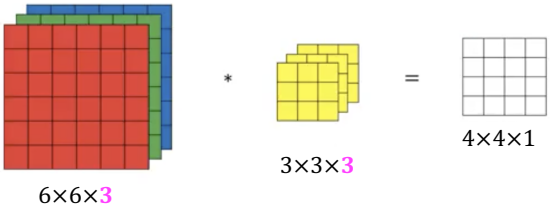

27 operations each time (28 if bias term)

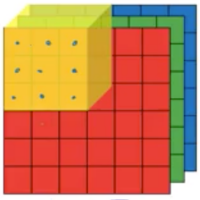

Each slice of the filter can detect different features in each of the 3 colour channels, for instance vertical edges in Red, horizontal edges in Green and vertical edges in Blue!

Andrew Ng https://www.youtube.com/watch?v=KTB_OFoAQcc&list=PLkDaE6sCZn6Gl29AoE31iwdVwSG-KnDzF&t=0s&index=7

19

Imperial College London

If Different Filters are Applied …

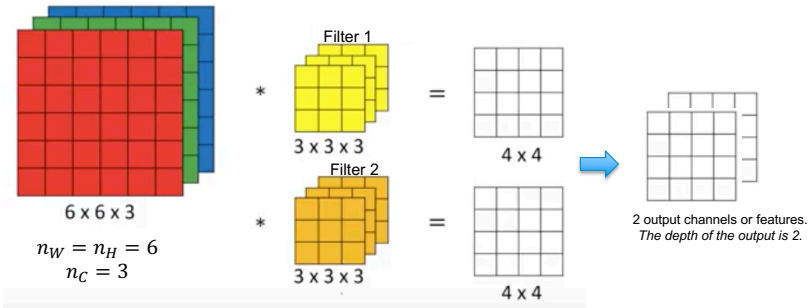

Exercise: *if my input Image is 50x50 with 10 channels, and if I apply 25 filters each of size 3x3x10 with a stride of 1, no padding, what is the size and number of channels of the output image?*

Answer: *size is 48x48. with 25 channels*

Andrew Ng https://www.youtube.com/watch?v=KTB_OFoAQcc&list=PLkDaE6sCZn6Gl29AoE31iwdVwSG-KnDzF&t=0s&index=7

20

Imperial College London

Number of Weights to Train for a Given Layer of the CNN

If I have 100 filters of size 4 x 4 x 3 (each with bias term) in layer l, what is the total number of weights (or degrees of freedom) to train for this layer?

$$\boldsymbol{Number\ of\ filters} \times (\boldsymbol{size\ of\ filter} + \mathbf{1}\ \boldsymbol{bias\ term}) =$$

$$\mathbf{100} \times (\mathbf{4} \times \mathbf{4} \times \mathbf{3} + \mathbf{1}) = \ \mathbf{100} \times \mathbf{49} = \mathbf{4900}\ \boldsymbol{weights}$$

Note that the number of parameters is independent of the size of the input image, which was not the case for feed-forward neural networks!

21

Imperial College London

Our first CNN Example: Recognize Whether Image of Cat or Dog

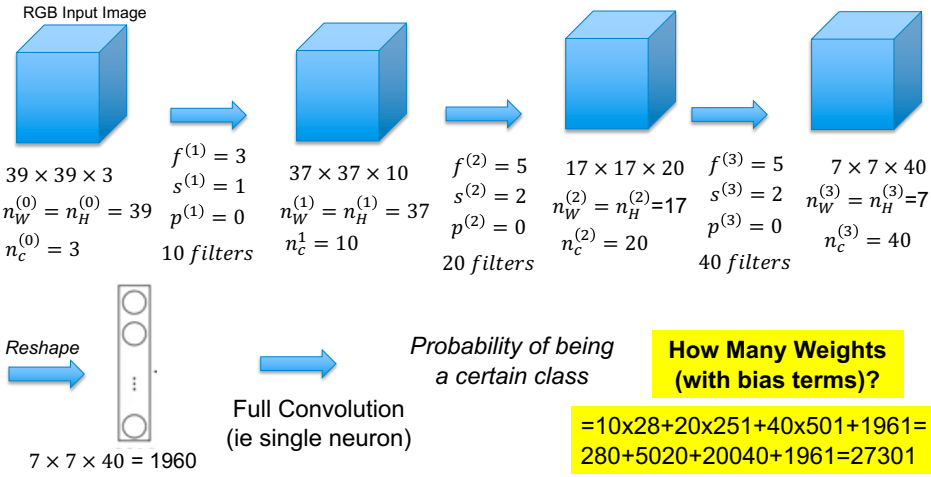

How Many Weights (with bias terms)?

=10x28+20x251+40x501+1961=
280+5020+20040+1961=27301

Inspired from Andre Ng

22

Imperial College London

Counting the Degrees of Freedom (or Parameters) for a layer l

Assume layer l has $n_c^{(l)}$ filters of size $f^{(l)}$, a padding $p^{(l)}$ and a stride $s^{(l)}$

Assume the input image has size $n_H^{(l-1)} \times n_W^{(l-1)} \times n_c^{(l-1)}$
and the output image has size $n_H^{(l)} \times n_W^{(l)} \times n_c^{(l)}$

The size of each filter is $f^{(l)} \times f^{(l)} \times n_c^{(l-1)}$ and the depth or number of output features is the number of filters $n_c^{(l)}$

And we can generalize the formula $n_H^{(l)} = \frac{n_H^{(l-1)} + 2p^{(l)} - f^{(l)}}{s^{(l)}} + 1$ and $n_W^{(l)} = \frac{n_W^{(l-1)} + 2p^{(l)} - f^{(l)}}{s^{(l)}} + 1$

And the total number of degrees of freedom (or parameters) of layer l is:

$$\boldsymbol{Number\ of\ filters} \times (\boldsymbol{size\ of\ filter} + \mathbf{1}\ \boldsymbol{bias\ term}) = \boldsymbol{n_c^{(l)}} \times \left(\boldsymbol{f^{(l)} \times f^{(l)} \times n_c^{(l-1)}} + \mathbf{1}\right)$$

23

Imperial College London

Convolutional Neural Networks

1. Convolutional Neural Networks
2. Pooling and Fully Connected Layers
3. Examples

24

Imperial College London

Three Types of Layers in a Convolutional Network

- Convolutional Layers (CONV)
- Fully Connected Layers (FC)
- Pooling Layers (POOL)

25

Imperial College London

Example of Max Pooling Layer (for Downsampling)

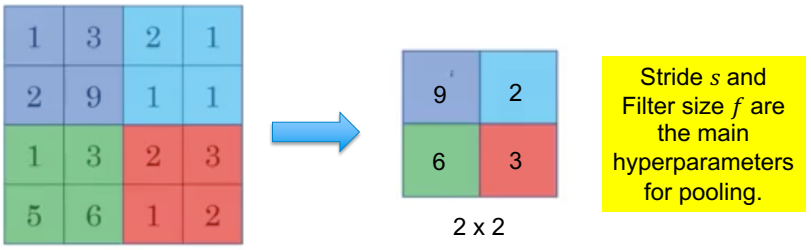

4 x 4

2 x 2

Stride s and Filter size f are the main hyperparameters for pooling.

Here $s = 2$ and $f = 2$ Usually $f = s$

Pooling makes the input representations (feature dimension) smaller and more manageable for the next layer, Max Pooling is used because it may be interesting to keep the high values for the activation of the next layer as they may characterize some important features. Pooling reduces the number of parameters and computations in the network, therefore controlling overfitting.

26

Imperial College London

We already know Fully Connected Layers!

Transform for example 32x32x3 image into 10x1 vector.

1. Reshape image into 1D vector x of dimension 32x32x3= 3072
2. Apply full convolution through matrix W

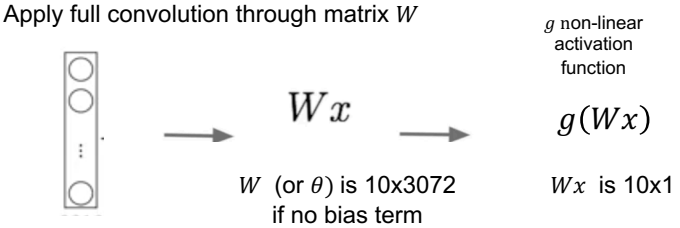

x is 3072x1

W (or θ) is 10x3072 if no bias term

Wx is 10x1

Fully connected layers often have many weights as they connect every neuron in input layer to every neuron in output layer.

27

Imperial College London

MNIST CNN Example (The LeNet-5 Architecture from LeCun)

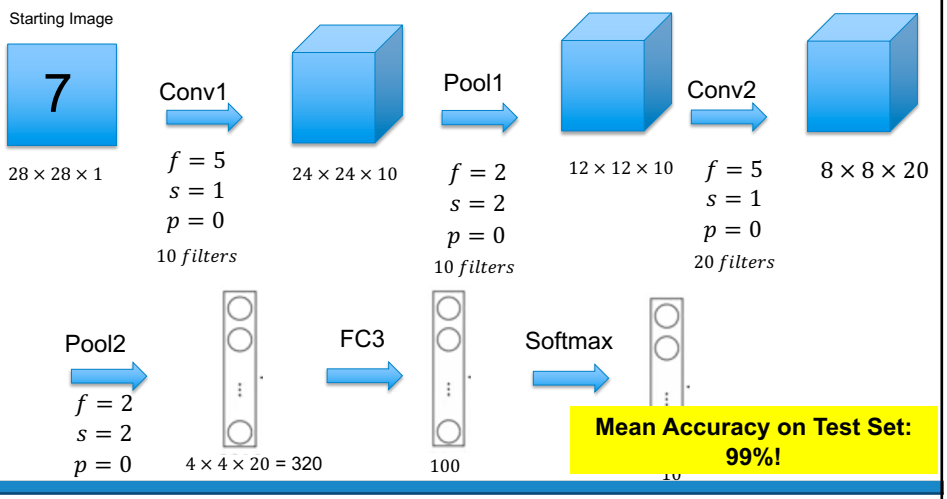

Mean Accuracy on Test Set: 99%!

28

Imperial College London

MNIST: Summary of CNN Layers (Convolutions with Bias Terms

	Size of input image n	Number of input channels	f	p	s	Size of output image (n+2p-f)/s+1	Number of output channels or filters	Number of output neurons	Size of Filter + 1	Number of Parameters
Conv1	28	1	5	0	1	24	10	5760	26	260
Pool1	24	10	2	0	2	12	10	1440		
Conv2	12	10	5	0	1	8	20	1280	251	5020
Pool2	8	20	2	0	2	4	20	320		
	Size of input							Number of output neurons		
FC1	320							100	321	32100
Softmax	100							10	101	1010
							Total Neurons	8910	Total Parameters	38390

Formula used for last column:
Number of Parameters= Number of output channels × (Size of filter+1) (as we assume there is a bias term)

29

Imperial College London

MNIST: An Interactive CNN Example

From Paper:
An Interactive Node-Link Visualization of Convolutional Neural Networks, by A. Harley

http://scs.ryerson.ca/~aharley/vis/conv/flat.html

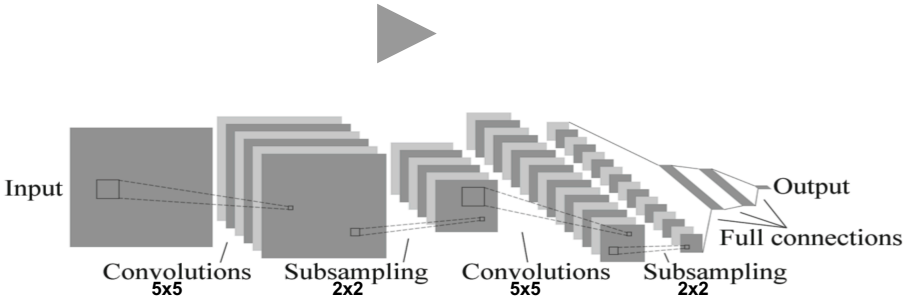

30

Imperial College London

Structure of Harley's Network

From Paper:
An Interactive Node-Link Visualization of Convolutional Neural Networks, by A. Hartey

	Size of input image n	Number of input channels	f	p	s	Size of output image (n+2p-f)/s+1	Number of output channels or filters	Number of output neurons	Size of Filter + 1	Number of Parameters
Conv1	32	1	5	0	1	28	6	4704	26	156
MaxPool	28	6	2	0	2	14	6	1176		
Conv2	14	6	5	0	1	10	16	1600	151	2416
MaxPool	10	16	2	0	2	5	128	3200		
	Size of input							Number of output neurons		
FC1	3200							120	3201	384120
FC2	120							100	121	12100
Softmax	100							10	101	1010
							Total Neurons	10910	Total Parameters	399802

Formula used for last column:
Number of Parameters= Number of output channels × (Size of filter+1) (as we assume there is a bias term)

31

Imperial College London

CNN Parameters are optimized the usual way

Example of cat image identification. Training set $(x^{(1)}, y^{(1)}) \ldots (x^{(m)}, y^{(m)})$

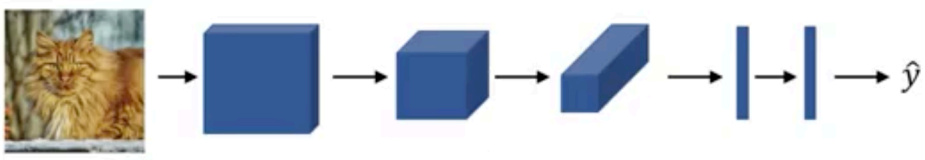

$Cost\ Function\ \mathrm{J} = \frac{1}{m}\sum_{i=1}^{m} \mathcal{L}(\hat{y}_i, y_i)$ (mean of cost functions for each data point)

The value of m depends whether batch, minibatch or stochastic gradient descent is used

32

Imperial College London

Convolutional Neural Networks

1. Convolutional Neural Networks
2. Pooling and Fully Connected Layers
3. Examples

33

Imperial College London

Some of the Best-Known Datasets

Very deep neural networks work best when trained on very large datasets!

- **MNIST: Handwritten digits, 60000 Training Images, 10000 Test Images**
- **CIFAR-10 / CIFAR-100: 50k Training, 10k Test Images of 10 (CIFAR-10) or 100 (CIFAR-100) classes**

  Color Images are 32x32, Task: Classification https://www.cs.toronto.edu/~kriz/cifar.html

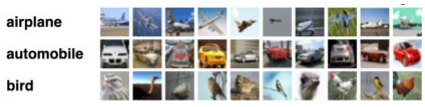

- **Imagenet: > 15 Million Images in 20,000 classes! https://en.wikipedia.org/wiki/ImageNet**

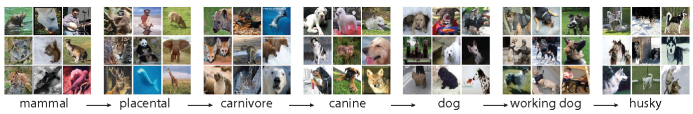

34

Imperial College London

Some of the best-known CNNs

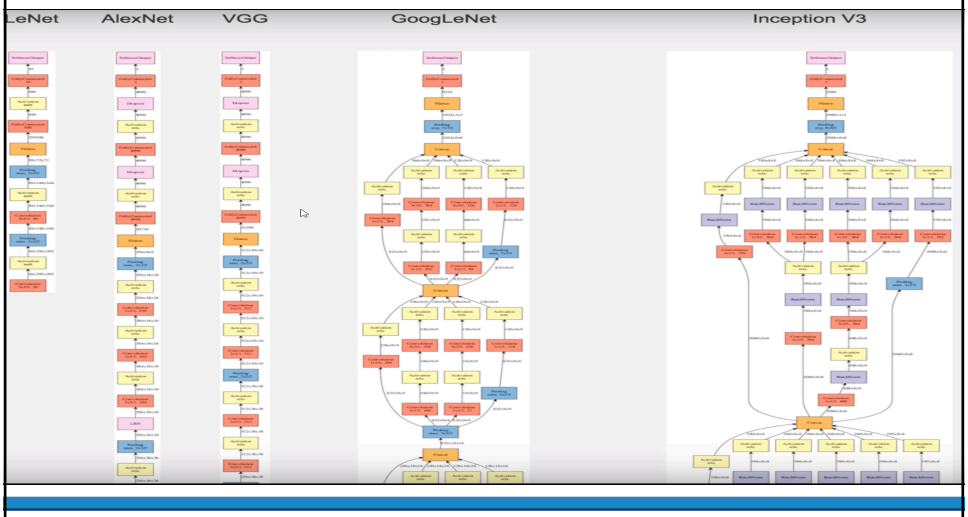

35

Imperial College London

The Image Large-Scale Visual Recognition Challenge

From 2010 to 2017, ImageNet has been running an annual competition (ILSVRC: Image Large Scale Visual Recognition Challenge) in visual recognition where participants are presented with the challenge of classifying 1.2 million images belonging to 1000 different classes from Imagenet dataset

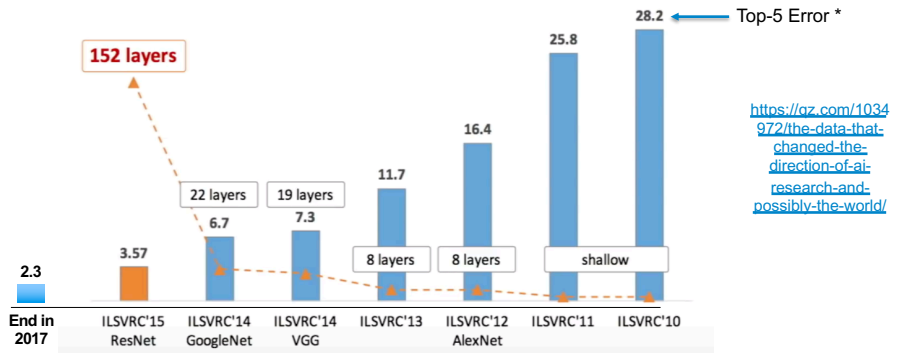

https://qz.com/1034972/the-data-that-changed-the-direction-of-ai-research-and-possibly-the-world/

** The **Top-5 error** is the percentage of the time that the classifier did not include the correct class among its **top 5** guesses.*

https://www.youtube.com/watch?v=DAOcjicFr1Y

36

Imperial College London

Some of the best-known CNNs

- **LeNet-5** (Le Cun et al, 1998) for Digit Recognition: Architecture is CONV-POOL-CONV-POOL-FC-FC
- **AlexNet** (ILSRVC 2012): Very similar to Le-Net-5, with just more layers (ReLU, Dropout 0.5 and Data Augmentation). About 60 MM parameters for 8 layers. First CNN-based winner of ILSRVC!
- **VGG** (co-winner of ILSRVC 2014): More layers (16 to 19) but very small filters for convolution (3x3, Pad 1, Stride 1, in stacked layers) and pooling (2x2, stride 2) . More non-linearities , with less parameters in CONV layers. But VGG16 has 138 MM parameters, mainly due to last Fully-Connected layers. VGG19 only slightly better than VGG16.

37

Imperial College London

Some of the best-known CNNs VGG

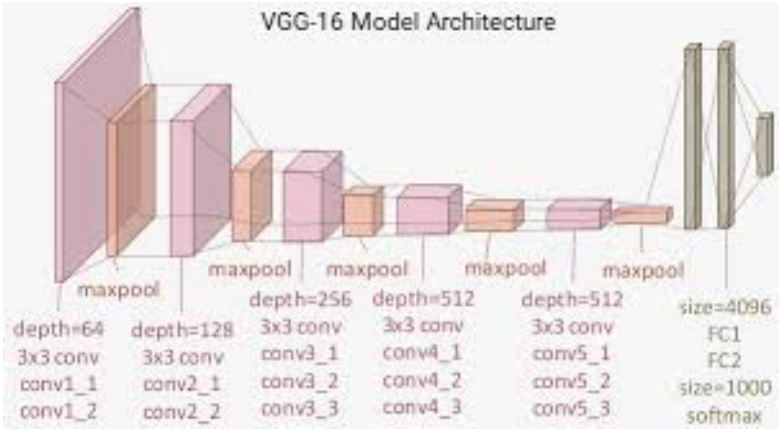

38

Imperial College London

Some of the best-known CNNs: Inception Module

- **GoogLeNet/Inception** (co-winner of ILSRVC 2014): More layers (22), based on "Inception" module ("network within a network"), no Fully-Connected layer, only 5 Million parameters! General Philosophy based on basic Inception Module ("Network within a Network")

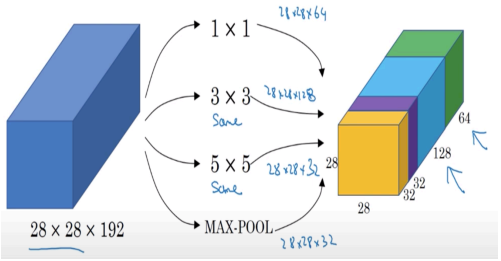

Example from Andrew Ng's C4W2L06

39

Imperial College London

Some of the best-known CNNs: ResNet

- **Resnet (winner of ILSRVC 2015):** Very deep network **(**152 layers) using residual connections. No FC layers. Did better than human performance!

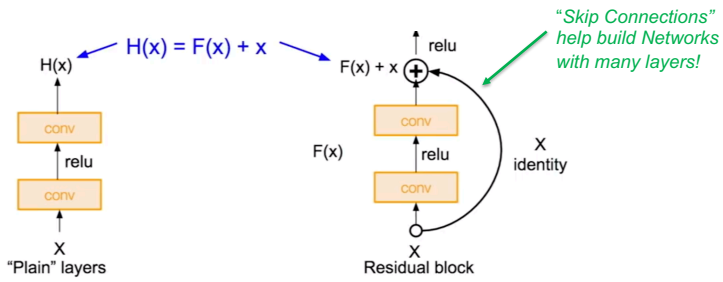

40

Imperial College London

Comparing Accuracy, Speed and Memory Usage

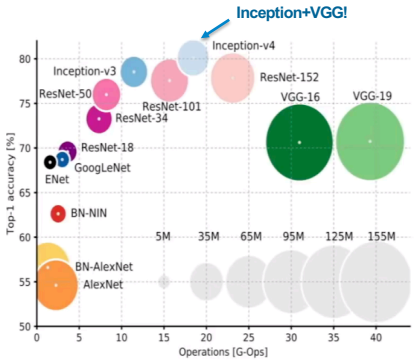

- **Download open model implementations pre-trained on large datasets:**

  Pytorch: Torchvision https://pytorch.org/docs/stable/torchvision/models.htm

https://www.youtube.com/watch?v=DAOcjicFr1Y

Imperial College London

Transfer Learning (1)

The most well-known CNN designs are available on-line and have been successfully trained on very large number of images (1,000,000s).

In many applications we often work from a relatively small number of images.

Why not start from an existing trained CNN sharing an objective of a similar nature (such as Classification) and tailor it to our application?

Imperial College London

Example of Transfer Learning Approach

Take existing trained network such as Inception v3 model, trained on ImageNet dataset for differentiating between 1,000 different classes of images.

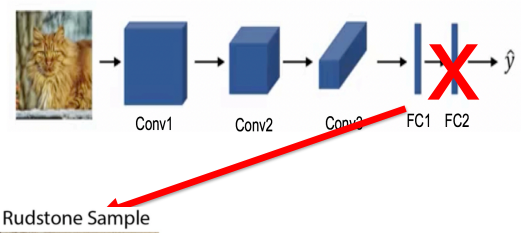

Rudstone Sample

Re-train last layer of the network on images of interest, such as pictures of carbonate cores .

From Sharinia Kanagandran and Cedric John, Imperial College

43

Imperial College London

Semantic Segmentation with Down- and Up-Sampling

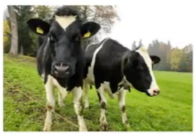

Input:
3 x H x W

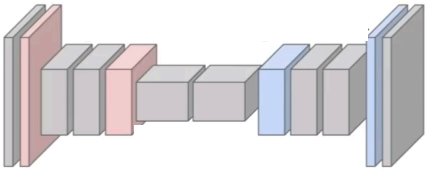

Predictions:
H x W

https://www.youtube.com/watch?v=azM57JuQpQI

44

Imperial College London

Segmentation Model: U-Net (Ronneberger et al, 2015)

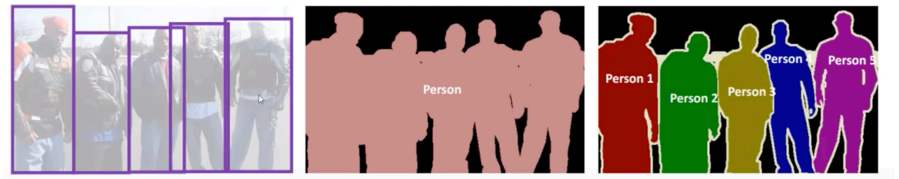

Object Detection | Semantic Segmentation | **Instance Segmentation**

https://www.youtube.com/watch?v=azM57JuQpQI

45

Imperial College London

Segmentation Model: U-Net (Ronneberger et al, 2015)

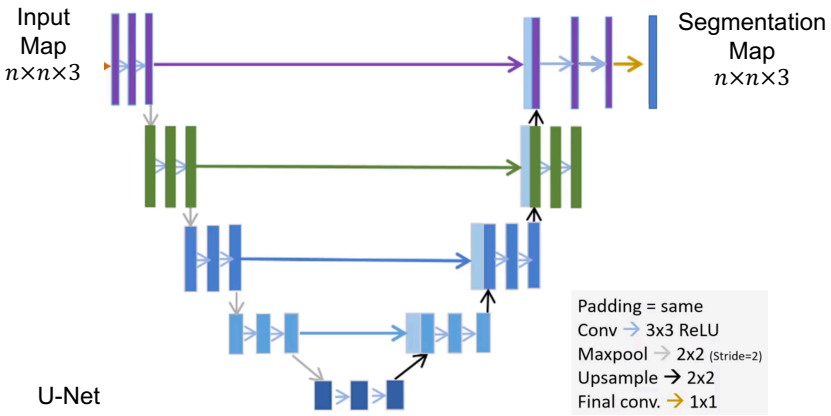

https://www.youtube.com/watch?v=azM57JuQpQI

46

Imperial College London

Applying U-Net to Seismic Interpretation

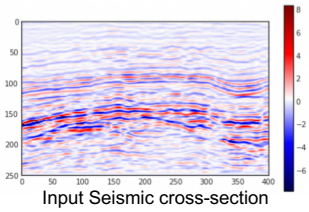

Input Seismic cross-section

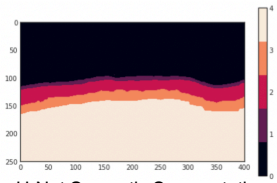

U-Net Semantic Segmentation

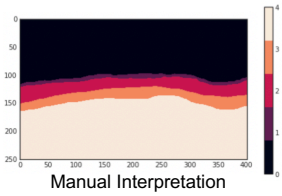

Manual Interpretation

47

Imperial College London

Conclusion

- Convolutional Neural Networks mostly used on 2-D images or 3-D Volumes for image classification, image segmentation, object localization/detection.
- Convolutional Neural networks are the basis for many successful Deep Learning Applications.
- AlexNet, VGG, GoogleNet/Inception and Resnet winners of ILSVRCs
- Transfer Learning: many complex but successful CNN architectures can be transferred from one application to another , by re-training the last layer.

48

Imperial College London

Coursework - To be Returned by Monday 9 am.

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky University of Toronto kriz@cs.utoronto.ca
Ilya Sutskever University of Toronto ilya@cs.utoronto.ca
Geoffrey E. Hinton University of Toronto hinton@cs.utoronto.ca

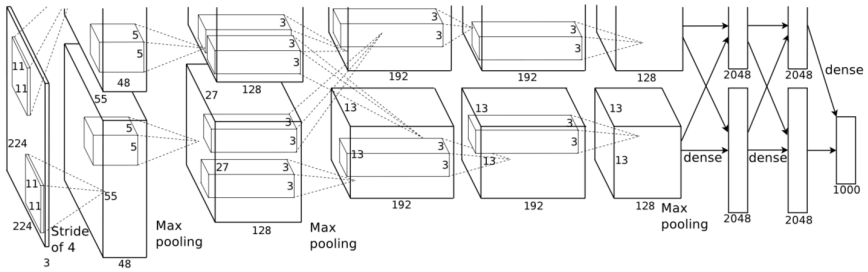

49

Imperial College London

Coursework

	Size of input image n	Number of input channels	f	p	s	Size of output image (n+2p-f)/s+1	Number of output channels or filters	Number of output neurons	Size of Filter + 1	Number of Parameters
Conv1	227	3	11	0	4		48			
MaxPool			3	0	2		48			
Conv2			6	2	1		128			
MaxPool			4	0	2		128			
Conv3			5	2	1		192			
Conv4			4	1	1		192			
Conv5			3	1	1		128			
MaxPool			4	0	3		128			
	Size of input							Number of output neurons		
FC1								1600		
FC2								1600		
Softmax								1000		
							Total Neurons		Total Parameters	

Formula used for last column has been given in the morning course:
Number of Parameters= Number of output channels × (Size of filter+1) (as we assume there is a bias term)

50

Imperial College London

Coursework

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca

Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

Delving Deep into Rectifiers:
Surpassing Human-Level Performance on ImageNet Classification

Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun

Microsoft Research
{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

Understanding the difficulty of training deep feedforward neural networks

Xavier Glorot **Yoshua Bengio**
DIRO, Université de Montréal, Montréal, Québec, Canada